

11785 HW3P2

Name: Xinkai Chen

AndrewID: xinkaic

Code: 11785_hw3p2_xinkaic.ipynb

Steps:

1. Download data and unzip them

2. Import libraries and run on GPU

3. Convert images to tensors, define dataloaders and load training data and dev data

- Dataset is pretty straightforward; in the `collate_fn` I padded inputs and labels together. I used `batch_first` on labels because they will be passed to `CTCLoss`.

- Dataloader has batch size of 64

4. Define network

- I tried different things:

- Baseline (as defined in Piazza): that gave me about a 11-12 distance

- Added CNN layer and extra Linear layer, also expanded the LSTM layers: I read a paper that used Conv – LSTM – FC, so I used a `Conv1d` layer with a small kernel size to preserve more features; I made the LSTM layers wider by using 512 hidden units; also deeper by using 4 stacked bi-LSTM layers; added another Linear layer before the output layer.

- I tried using more Conv layers but that didn't work well, so I only kept one layer. I discussed with classmates after the deadline, it turns out that using pooling/downsampling can be helpful.

- I didn't try some regularization tricks on LSTM when I reached distance < 8; but also during the discussion, Dropout can be helpful, and I will definitely look into those when doing HW4.

5. Define training parameters

- Criterion: `CTCLoss()`, I struggled with `batch_first` at first but then managed correct it; I read the documentation carefully to send in the correct inputs.

- Learning rate: starting with 0.001, *0.85 after each epoch (I made `step_size` larger when the training was just started, and then decrease learning rate more frequently as the loss decreases more)

- Optimizer: Adam with weight decay of $5e-6$ and the learning rate above

6. Training

- The distance reaches the A cut-off after around 25 epochs

- Validate after each epoch

7. Evaluate results and tune hyperparameters

- I tried different networks and tried different hyperparameters. And then I evaluate the model on the validation set. Also I print out the outputs to see if my model generation makes sense.

8. When the val loss & distance is okay, load the test set and predict on test set.

9. Generate submission.csv and submit to Kaggle.

